

Loading and Preprocessing Data

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV, KFold, train_test_split
from sklearn.utils.class_weight import compute_class_weight
from sklearn import metrics
from sklearn.metrics import roc_curve, roc_auc_score

# Ignore warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
data = pd.read_csv('/content/DATA File.csv')
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    569 non-null    int64
1   Diagnosis             569 non-null    object
2   radius1               569 non-null    float64
3   texture1              569 non-null    float64
4   perimeter1            569 non-null    float64
5   area1                 569 non-null    float64
6   smoothness1           569 non-null    float64
7   compactness1          569 non-null    float64
8   concavity1            569 non-null    float64
9   concave_points1       569 non-null    float64
10  symmetry1             569 non-null    float64
11  fractal_dimension1    569 non-null    float64
12  radius2               569 non-null    float64
13  texture2              569 non-null    float64
14  perimeter2            569 non-null    float64
15  area2                 569 non-null    float64
16  smoothness2           569 non-null    float64
17  compactness2          569 non-null    float64
18  concavity2            569 non-null    float64
19  concave_points2       569 non-null    float64
20  symmetry2             569 non-null    float64
21  fractal_dimension2    569 non-null    float64
22  radius3               569 non-null    float64
23  texture3              569 non-null    float64
24  perimeter3            569 non-null    float64
25  area3                 569 non-null    float64
26  smoothness3           569 non-null    float64
27  compactness3          569 non-null    float64
28  concavity3            569 non-null    float64
29  concave_points3       569 non-null    float64
30  symmetry3             569 non-null    float64
31  fractal_dimension3    569 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB
```

data

	ID	Diagnosis	radius1	texture1	perimeter1	area1	smoothness1	compactness1	concav:
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.00
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.10
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.20
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.10
...	...	...	...	...	...	...	...	...	...
564	926424	M	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.20
565	926682	M	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.10
566	926954	M	16.60	28.08	108.30	858.1	0.08455	0.10230	0.00
567	927241	M	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.30
568	92751	B	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00

569 rows × 32 columns


```
# Counting missing values
missing_values = data.isnull().sum()
missing_values
```

```
ID          0
Diagnosis    0
radius1      0
texture1     0
perimeter1   0
area1        0
smoothness1  0
compactness1 0
concavity1   0
concave_points1 0
symmetry1    0
fractal_dimension1 0
radius2      0
```

```
texture2      0
perimeter2    0
area2         0
smoothness2   0
compactness2  0
concavity2    0
concave_points2 0
symmetry2     0
fractal_dimension2 0
radius3       0
texture3      0
perimeter3    0
area3         0
smoothness3   0
compactness3  0
concavity3    0
concave_points3 0
symmetry3     0
fractal_dimension3 0
dtype: int64
```

```
# Identify duplicate rows
duplicates = data.duplicated()
print("Duplicate Rows:\n", data[duplicates])
```

```
# Count the number of duplicate rows
num_duplicates = duplicates.sum()
print(f"Number of Duplicate Rows: {num_duplicates}")
```




Duplicate Rows:  
Empty DataFrame  
Columns: [ID, Diagnosis, radius1, texture1, perimeter1, area1, smoothness1, compactness1, concavity1, concave\_points1, symmetry1, fractal\_dimension1, r  
Index: []

[0 rows x 32 columns]

Number of Duplicate Rows: 0

```
# Drop 'ID' column
new_data = data.drop(['ID'], axis=1)
```

```
print(new_data.describe())
```



	radius1	texture1	perimeter1	area1	smoothness1	\
count	569.000000	569.000000	569.000000	569.000000	569.000000	
mean	14.127292	19.289649	91.969033	654.889104	0.096360	
std	3.524049	4.301036	24.298981	351.914129	0.014064	
min	6.981000	9.710000	43.790000	143.500000	0.052630	
25%	11.700000	16.170000	75.170000	420.300000	0.086370	
50%	13.370000	18.840000	86.240000	551.100000	0.095870	
75%	15.780000	21.800000	104.100000	782.700000	0.105300	
max	28.110000	39.280000	188.500000	2501.000000	0.163400	

	compactness1	concavity1	concave_points1	symmetry1	\
count	569.000000	569.000000	569.000000	569.000000	
mean	0.104341	0.088799	0.048919	0.181162	
std	0.052813	0.079720	0.038803	0.027414	
min	0.019380	0.000000	0.000000	0.106000	
25%	0.064920	0.029560	0.020310	0.161900	
50%	0.092630	0.061540	0.033500	0.179200	
75%	0.130400	0.130700	0.074000	0.195700	
max	0.345400	0.426800	0.201200	0.304000	

	fractal_dimension1	...	radius3	texture3	perimeter3	\
count	569.000000	...	569.000000	569.000000	569.000000	
mean	0.062798	...	16.269190	25.677223	107.261213	
std	0.007060	...	4.833242	6.146258	33.602542	
min	0.049960	...	7.930000	12.020000	50.410000	
25%	0.057700	...	13.010000	21.080000	84.110000	
50%	0.061540	...	14.970000	25.410000	97.660000	
75%	0.066120	...	18.790000	29.720000	125.400000	
max	0.097440	...	36.040000	49.540000	251.200000	

	area3	smoothness3	compactness3	concavity3	concave_points3	\
count	569.000000	569.000000	569.000000	569.000000	569.000000	
mean	880.583128	0.132369	0.254265	0.272188	0.114606	
std	569.356993	0.022832	0.157336	0.208624	0.065732	
min	185.200000	0.071170	0.027290	0.000000	0.000000	
25%	515.300000	0.116600	0.147200	0.114500	0.064930	
50%	686.500000	0.131300	0.211900	0.226700	0.099930	
75%	1084.000000	0.146000	0.339100	0.382900	0.161400	
max	4254.000000	0.222600	1.058000	1.252000	0.291000	

	symmetry3	fractal_dimension3
count	569.000000	569.000000
mean	0.290076	0.083946
std	0.061867	0.018061
min	0.156500	0.055040
25%	0.250400	0.071460
50%	0.282200	0.080040
75%	0.317900	0.092080
max	0.663800	0.207500

[8 rows x 30 columns]

## ✕ Exploratory Data Analysis

```
# Visualization of each target and features
import math

# Number of features
num_features = len(new_data.columns)

# Calculate the number of rows and columns for subplots
n_cols = 3
n_rows = math.ceil(num_features / n_cols)

# Create subplots
fig, axes = plt.subplots(n_rows, n_cols, figsize=(15, 5 * n_rows))
```















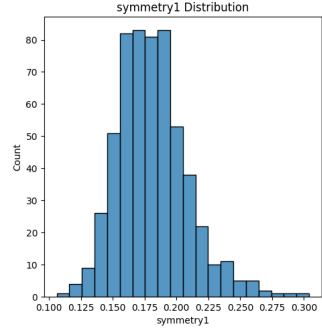









```
fig, axes = plt.subplots(n_rows, n_cols, figsize=(10, 5 * n_cols))


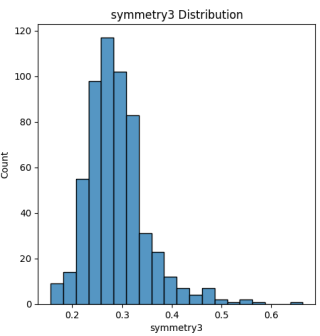





# Flatten axes array for easy iteration
axes = axes.flatten()

# Loop over each column to create histograms
for i, col in enumerate(new_data.columns):
    sns.histplot(data[col], bins=20, ax=axes[i])
    axes[i].set_title(f'{col} Distribution')

# Hide any unused subplots
for j in range(i + 1, n_rows * n_cols):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```





```
df = pd.get_dummies(new_data, columns=['Diagnosis'], dtype=int)
```

```
# Checking the count of the instances for each class in the target variable.
class_counts = new_data['Diagnosis'].value_counts()
print(class_counts)
```

```
➡ Diagnosis
B      357
M      212
Name: count, dtype: int64
```

```
# Checking the count of the instances for each class in the target variable.
class_counts = new_data['Diagnosis'].value_counts()
print("Class counts:\n", class_counts)
```

```
# Calculate the ratio of the minority class to the majority class
ratio = class_counts.min() / class_counts.max()
print(f"\nRatio of minority to majority class: {ratio:.2f}")
```

```
# Determine if oversampling is needed
if ratio < 0.5:
    print("The dataset is imbalanced and may benefit from oversampling.")
else:
    print("The dataset is relatively balanced and may not need oversampling.")
```

```
➡ Class counts:
Diagnosis
B      357
M      212
Name: count, dtype: int64

Ratio of minority to majority class: 0.59
The dataset is relatively balanced and may not need oversampling.
```

```
X = new_data.drop('Diagnosis', axis=1)
y = new_data['Diagnosis']
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=44)
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import KFold
```

## Model Training and Evaluation

```
import sklearn.metrics as metrics

def print_model_performance(model, X_test, y_test):
    y_pred = model.predict(X_test)
    print('Accuracy score:\n', metrics.accuracy_score(y_test, y_pred))
    print('Confussion matrix:\n',metrics.confusion_matrix(y_test, y_pred))
    print('Classification report:\n', metrics.classification_report(y_test, y_pred, digits=4))
    # Set pos_label to one of your actual labels
    print('Precision:\n',metrics.precision_score(y_test,y_pred, pos_label='M')) # Change 'M' if needed
    print('Recall:\n', metrics.recall_score(y_test, y_pred, pos_label='M'))    # Change 'M' if needed
```

## Logistic Regression Classifier & It's ROC Curve

```
import sklearn.metrics as metrics

lr_classifier = LogisticRegression(max_iter=1000)
desc = lr_classifier.fit(X_train , y_train)
print(desc)

y_pred = lr_classifier.predict(X_test)
print(metrics.confusion_matrix(y_test, y_pred))
print(metrics.classification_report(y_test , y_pred))
```

```
LogisticRegression(max_iter=1000)
[[74  1]
 [ 2 37]]
precision    recall  f1-score   support

      B       0.97       0.99       0.98        75
      M       0.97       0.95       0.96        39

 accuracy          0.97          114
 macro avg       0.97       0.97       0.97          114
weighted avg       0.97       0.97       0.97          114
```

```
from sklearn.utils.class_weight import compute_class_weight

# Find all unique labels in y_train
unique_labels = y_train.unique()

# Calculate class weights using the unique labels
class_weights = compute_class_weight('balanced', classes=unique_labels, y=y_train)

# Create a logistic regression classifier with class weights
lr_classifier = LogisticRegression(class_weight=dict(zip(unique_labels, class_weights)))

# Train and evaluate the model
lr_classifier.fit(X_train, y_train)
y_pred = lr_classifier.predict(X_test)
print_model_performance(lr_classifier, X_test, y_test)
```

```
Accuracy score:
0.9824561403508771
Confussion matrix:
[[75  0]
 [ 2 37]]
Classification report:
precision    recall  f1-score   support

      B       0.9740       1.0000       0.9868        75
      M       1.0000       0.9487       0.9737        39

 accuracy          0.9825          114
 macro avg       0.9870       0.9744       0.9803          114
weighted avg       0.9829       0.9825       0.9823          114


Precision:
1.0
Recall:
0.9487179487179487
```

```
from sklearn.metrics import roc_curve, roc_auc_score

# Generate ROC curve
# Calculate the predicted probabilities for the positive class (e.g., 'M')
y_pred_proba = lr_classifier.predict_proba(X_test)[: , 1]

# Compute ROC curve and ROC AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba, pos_label='M')
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, linestyle='--', color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.savefig('ROC',dpi=300)
plt.show()
```



## Logistic Regression Classifier with Hyperparameter Tuning

```
# Define the logistic regression classifier
lr_classifier = LogisticRegression()

# Define the values you want to search for C, penalty, and solver
lr_C = np.array([50])
lr_penalty = ['l1', 'elasticnet'] # Changed 'l2' to 'l1' and 'elasticnet'
lr_solver = ['lbfgs', 'liblinear', 'sag'] # Added more solver options

grid_values = {'C': lr_C, 'penalty': lr_penalty, 'solver': lr_solver} # L1 = Lasso L2 = Ridge

# Specify the cross-validation strategy
cross_validation = KFold(n_splits=5, shuffle=False)

# Create the GridSearchCV object
lrc_grid = GridSearchCV(lr_classifier, param_grid=grid_values, cv=cross_validation, n_jobs=-1, scoring='f1_macro')

# Fit the grid search to your training data
desc = lrc_grid.fit(X_train, y_train)

# Print the best parameters and model performance
print("\n The best parameters across all searched params:\n", lrc_grid.best_params_)
print_model_performance(lrc_grid, X_test, y_test) # You should have a print_model_performance function defined elsewhere
```

```
The best parameters across all searched params:
{'C': 50, 'penalty': 'l1', 'solver': 'liblinear'}
Accuracy score:
0.9912280701754386
Confussion matrix:
[[75  0]
 [ 1 38]]
Classification report:
              precision    recall  f1-score   support

      B       0.9868       1.0000       0.9934         75
      M       1.0000       0.9744       0.9870         39

   accuracy       0.9934       0.9872       0.9912        114
  macro avg       0.9934       0.9872       0.9902        114
weighted avg       0.9913       0.9912       0.9912        114

Precision:
1.0
Recall:
0.9743589743589743
```

## RandomForest Classifier & It's ROC Curve

```
from sklearn.ensemble import RandomForestClassifier

# Define and train a Random Forest classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=44)
rf_classifier.fit(X_train, y_train)

# Evaluate the model
y_pred = rf_classifier.predict(X_test)
print_model_performance(rf_classifier, X_test, y_test) # Define a print_model_performance function
```

```
Accuracy score:
0.9649122807017544
Confussion matrix:
[[73  2]
 [ 2 37]]
Classification report:
              precision    recall  f1-score   support

      B       0.9733       0.9733       0.9733         75
      M       0.9487       0.9487       0.9487         39

   accuracy       0.9649       0.9610       0.9649        114
  macro avg       0.9610       0.9610       0.9610        114
weighted avg       0.9649       0.9649       0.9649        114


Precision:
0.9487179487179487
Recall:
```



0.9487179487179487

```
# Generate the ROC curve
y_pred_proba = rf_classifier.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba, pos_label='M') # Specify the positive class
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, linestyle='--', color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.savefig('ROC',dpi=300)
plt.show()
```



## Support Vector Machine (SVM) Classifier & It's ROC Curve

```
from sklearn import svm
svm_classifier = svm.SVC() #default kernel is radial basic function (RBF) kernel

desc = svm_classifier.fit(X_train, y_train)
print(desc)
print_model_performance(svm_classifier, X_test, y_test)
```

```
SVC()
Accuracy score:
  0.956140350877193
Confussion matrix:
[[73  2]
 [ 3 36]]
Classification report:

```

	precision	recall	f1-score	support
B	0.9605	0.9733	0.9669	75
M	0.9474	0.9231	0.9351	39
accuracy			0.9561	114
macro avg	0.9539	0.9482	0.9510	114
weighted avg	0.9560	0.9561	0.9560	114

```


Precision:
 0.9473684210526315
Recall:
 0.9230769230769231
```

```
import sklearn.metrics as metrics

for C in [1,10,100,1000,10000,100000,1000000]:
    for kernel in ['rbf','poly']:
        svm_classifier = svm.SVC(C=C, kernel=kernel) #similarly.gamma=0.01, 0.5
        svm_classifier.fit(X_train, y_train)
        y_pred = svm_classifier.predict(X_test)
        scores = metrics.accuracy_score(y_test, y_pred)
        print(f'C={C:<10} kernel={kernel:<10} accuracy score={scores: .3%}')
```

```
C=1          kernel=rbf          accuracy score= 95.614%
C=1          kernel=poly         accuracy score= 94.737%
C=10         kernel=rbf          accuracy score= 97.368%
C=10         kernel=poly         accuracy score= 96.491%
C=100        kernel=rbf          accuracy score= 97.368%
C=100        kernel=poly         accuracy score= 96.491%
C=1000       kernel=rbf          accuracy score= 97.368%
C=1000       kernel=poly         accuracy score= 96.491%
C=10000      kernel=rbf          accuracy score= 97.368%
C=10000      kernel=poly         accuracy score= 98.246%
C=100000     kernel=rbf          accuracy score= 98.246%
C=100000     kernel=poly         accuracy score= 96.491%
C=1000000    kernel=poly         accuracy score= 97.368%
```

```
sns.heatmap(metrics.confusion_matrix(y_test, y_pred),annot=True, fmt='d', cmap='rainbow')
```



```
# Define the SVM classifier
svm_classifier = svm.SVC(probability=True) # Set probability=True to enable predict_proba method

# Define the parameter grid
param_grid = {
    'C': [0.1, 1, 10, 100],
    'gamma': [1, 0.1, 0.01, 0.001],
    'kernel': ['rbf']
}

# Initialize GridSearchCV with SVM classifier and the parameter grid
grid = GridSearchCV(svm_classifier, param_grid, refit=True, verbose=2, cv=3)

# Fit the GridSearchCV to the data
grid.fit(X_train, y_train.values.ravel())

# Print the best parameters and estimator
print(f"Best parameters: {grid.best_params_}")
print(f"Best estimator: {grid.best_estimator_}")

# Evaluate the best model found by GridSearchCV
best_model = grid.best_estimator_
print_model_performance(best_model, X_test, y_test)

# Generate the ROC curve for the best model
y_pred_proba = best_model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba, pos_label='M') # Specify 'M' as the positive class
roc_auc = roc_auc_score(y_test, y_pred_proba)
# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, linestyle='--', color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```




```
Fitting 3 folds for each of 16 candidates, totalling 48 fits
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=0.1, gamma=0.01, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=0.1, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=0.1, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=1, gamma=1, kernel=rbf; total time= 0.0s
[CV] END .....C=1, gamma=1, kernel=rbf; total time= 0.1s
[CV] END .....C=1, gamma=1, kernel=rbf; total time= 0.0s
[CV] END .....C=1, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=1, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=1, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=1, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=1, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=1, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=1, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=1, gamma=0.001, kernel=rbf; total time= 0.1s
[CV] END .....C=1, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=10, gamma=1, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=1, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=1, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=0.001, kernel=rbf; total time= 0.1s
[CV] END .....C=10, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=100, gamma=1, kernel=rbf; total time= 0.1s
[CV] END .....C=100, gamma=1, kernel=rbf; total time= 0.1s
[CV] END .....C=100, gamma=1, kernel=rbf; total time= 0.1s
[CV] END .....C=100, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=100, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=100, gamma=0.1, kernel=rbf; total time= 0.1s
[CV] END .....C=100, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=100, gamma=0.01, kernel=rbf; total time= 0.1s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=100, gamma=0.001, kernel=rbf; total time= 0.1s
```

Best parameters: {'C': 10, 'gamma': 0.001, 'kernel': 'rbf'}  
Best estimator: SVC(C=10, gamma=0.001, probability=True)


Accuracy score:  
0.9210526315789473  
Confussion matrix:  
[[67 8]  
[ 1 38]]

Classification report:					
	precision	recall	f1-score	support	
B	0.9853	0.8933	0.9371	75	
M	0.8261	0.9744	0.8941	39	
accuracy				0.9211	114
macro avg	0.9057	0.9338	0.9156	114	
weighted avg	0.9308	0.9211	0.9224	114	

Precision:  
0.8260869565217391  
Recall:  
0.9743589743589743



```
sns.heatmap(metrics.confusion_matrix(y_test, y_pred),annot=True, fmt='d', cmap='rainbow')
```



## ✕ K-Nearest Neighbors Classifier & It's ROC Curve

```
knn_classifier = KNeighborsClassifier(n_neighbors=4)
knn_classifier.fit(X_train, y_train)

desc = knn_classifier.fit(X_train, y_train)
print(desc)
print_model_performance(knn_classifier, X_test, y_test)
```

↗ KNeighborsClassifier(n\_neighbors=4)

Accuracy score:  
0.9649122807017544

Confussion matrix:  
[[73 2]  
[ 2 37]]

Classification report:

	precision	recall	f1-score	support
B	0.9733	0.9733	0.9733	75
M	0.9487	0.9487	0.9487	39
accuracy			0.9649	114
macro avg	0.9610	0.9610	0.9610	114
weighted avg	0.9649	0.9649	0.9649	114


Precision:  
0.9487179487179487

Recall:  
0.9487179487179487

```
# Predict probabilities
y_pred_proba = knn_classifier.predict_proba(X_test)[:, 1]

# Calculate ROC curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba, pos_label='M') # Specify 'M' as the positive class

# Plot ROC curve
plt.plot(fpr, tpr, linestyle='--', color='orange', label='KNN')
plt.title('ROC Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='best')
plt.savefig('ROC', dpi=300)
plt.show()
```



## ✕ Decision Tree Classifier & It's ROC Curve

```
from sklearn.tree import DecisionTreeClassifier
dt_classifier = DecisionTreeClassifier(criterion= "entropy" ,max_depth=10, random_state=44)
dt_classifier.fit(X_train,y_train)

print(dt_classifier)
print_model_performance(dt_classifier, X_test, y_test)
```

```
DecisionTreeClassifier(criterion='entropy', max_depth=10, random_state=44)
Accuracy score:
0.956140350877193
Confussion matrix:
[[72  3]
 [ 2 37]]
Classification report:

```

	precision	recall	f1-score	support
B	0.9730	0.9600	0.9664	75
M	0.9250	0.9487	0.9367	39
accuracy			0.9561	114
macro avg	0.9490	0.9544	0.9516	114
weighted avg	0.9566	0.9561	0.9563	114

## Hypermpparameter tuning with Desision Tree

```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, roc_curve, roc_auc_score, classification_report, confusion_matrix, precision_score, recall_score
from sklearn.tree import DecisionTreeClassifier
# Set up the parameter grid for the Decision Tree
param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Initialize a Decision Tree classifier
dt_classifier = DecisionTreeClassifier(random_state=44)
```